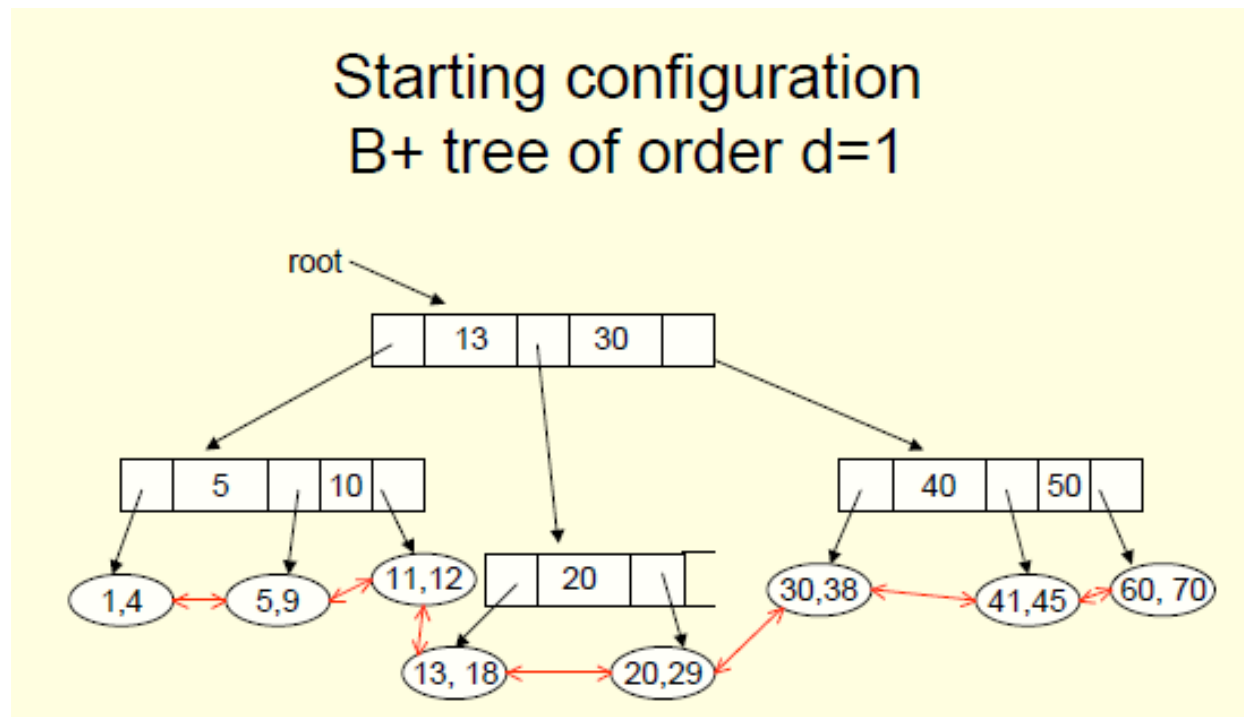



# Testing Mechanism for B+ Trees

- 以下幾個練習題，希望能夠幫助大家想想怎麼去測試你們所建構的 **B+ Trees**。
- 它們的資料結構跟你們的 **B+ Trees** 或許並不完全相同。
- 練習題純粹做為參考，並沒有任何實作的強迫性。

## Exercise 1:




Insert 19:  
split leaf; expand parent with key 18




# Testing B+ Trees (2)

Insert 27  
split leaf; expand parent with key 27 => too full




Insert 27  
 split leaf; split parent;  
 expand grandparent with key 20 => too full




# Testing B+ Trees (3)

Insert 27  
 split leaf; split parent; split grandparent  
 new root with key 20




Delete 5, then 9  
redistribute from right sibling




# Testing B+ Trees (4)

Delete 12  
merge leaves, delete key from parent




Delete 4, then 11  
merge leaves, delete key from parent  
=>parent not full enough



# Testing B+ Trees (5)


Delete 4, then 11  
merge leaves, merge parent, bringing down key 13  
=>grandparent not full enough



# Testing B+ Trees (6)

Delete 4, then 11

merge leaves; merge parent, bringing down key 13  
merge grandparent, bring down key 20,  
remove root



# Testing B+ Trees (7)

## Exercise 2:

- Assume this problem deals with the creation and use of a B+ Tree of order  $d=1$ . The leaves hold a maximum of 2 data entries.

**Step 1.** Bulk load a B+ Tree with data entries with keys values 2, 4, 6, 8, 10, 12, 14, 16 so that each leaf is **full** (or **half-full**). Show the state of the tree.

**Step 2.** Insert a data entry with key value 5 in your B+ Tree of Step 1. Show the state of the tree after the execution of the insert algorithm.

**Step 3.** Delete the data entry with key value 12 from your B+ Tree of Step 1. Show the final tree.

# Testing B+ Trees (8)

## Exercise 3:

- Initialize a B+ tree of order 2. Make your program insert the following integers, in this order:  
2, 16, 7, 20, 29, 33, 39, 38, 3, 5, 14, 19, 22, 24, 27, 34
- Insert 23 into the B+ Tree. Show the state of the tree.
- Delete 29, 33, and 34 from the B+ Tree. Show the final tree.



# Testing B+ Trees (9)

- 以下定義適用於 **Exercise 4** 至 **Exercise 6** 。
- number of search keys =  $x$   $\rightarrow$  number of pointers  $p = x+1$
- number of pointers  $p \rightarrow$  search keys  $d = p-1$
- Each node that is not a root or a leaf has between  $p/2$  and  $p$  children, where  $p$  is fixed for a particular tree.
- A leaf node has between  $(p-1)/2$  and  $p-1$  values (search keys)
- Size of one node =  $2x + 1 = x + p$

# Testing B+ Trees (10)

## Exercise 4:

- Consider a relational table R consisting of 12 tuples with primary key values 40, 20, 10, 30, 15, 55, 50, 21, 3, 54, 2, 8. Perform the following operations:
- Index it with a B+ tree with  $x = 3$  (number of search keys) on the primary key attribute. Insert the tuples in the above-specified order. How does the resulting tree look like?
- Delete values 7, 21, 50 from the index.

# Testing B+ Trees (11)

## Exercise 5:

- Consider a relational table R consisting of the following set of key values: 2, 3, 5, 7, 11, 17, 19, 23, 31.

**Step 1.** Assume that the tree is initially empty and values are added in ascending order. Construct B+ tree for the cases where the number of pointers is as follows:

a) Four b) Six c) Eight

**Step 2.** For each B+ tree of Step 1, show the form of the tree after each of the following series of operations:

a) Insert 9. b) Insert 10. c) Insert 8.  
d) Delete 23. e) Delete 19.

# Testing B+ Trees (12)

## Exercise 6:

- For a B+ tree with maximum keys = 4, insert the following keys in order: 10, 20, 30, 40, 50, 60, 70, 80, 90.
- Assuming keys increasing by 10, what is the first key added that causes the B+ tree to grow to height 3?
- Show the tree after deleting the following keys: 70, 90, 10.

- In the following, you find three algorithms that might help you to implement the B+ Tree.

```
function containsValue(int nodeID, int value) : boolean  
    node  $\leftarrow$  get the node for nodeID;  
    if node is a leaf node then  
        | if value is found in node then  
        | | return true  
        | else  
        | | return false  
    else // node is a branch node, continue with child  
        | childID  $\leftarrow$  child ID with greatest key  $k$  so that  $k \leq$  value in node;  
        | return containsValue(childID, value)
```

**Algorithm 1:** Checking if a value is contained in a B+ Tree.

```

function insertValue (int nodeID, int value) : long
    node  $\leftarrow$  get the node for nodeID;
    if node is a leaf node then
        if node already contains value then
            | return NO_CHANGES
        else // value has to be inserted
            increment the number of values stored in the tree;
            if some space is left then
                | insert the value into node;
                | return NO_CHANGES
            else // node has to be split
                | rightID  $\leftarrow$  create a new leaf node;
                | right  $\leftarrow$  the node with ID rightID;
                | distribute values between node and right;
                | return keyIDPair (firstValue (right), rightID)
    else // node is a branch node
        childID  $\leftarrow$  child ID with greatest key  $k$  so that  $k \leq$  value in node;
        result  $\leftarrow$  insertValue (childID, value);
        if result = NO_CHANGES then // nothing to do
            | return NO_CHANGES
        else // child was split
            | midKey  $\leftarrow$  getMidKey (result);
            | rightID  $\leftarrow$  getChildID (result);
            | if some space is left in node then
            |     | insert the rightID into node's child ID list;
            |     | insert the midKey into node's value list;
            |     | return NO_CHANGES
            | else // node is full and has to be split
            |     | rightID  $\leftarrow$  create a new branch node;
            |     | right  $\leftarrow$  the node with ID rightID;
            |     | distribute values and child pointers in node and right;
            |     | midKey'  $\leftarrow$  middle value between those in node and right;
            |     | return keyIDPair (midKey', rightID)

```

**Algorithm 2:** Inserting a value into a B+ Tree.

```

function deleteValue (int nodeID, int value) : boolean
    node ← get the node for nodeID;
    if node is a leaf node then
        if value is not found in node then
            | return true
        else // value is found
            delete value from node;
            decrement the number of values stored in the tree;
            if node still has enough values stored then
                | return true
            else // node is under-full
                | return false
    else // node is a branch node
        childID ← child ID with greatest key  $k$  so that  $k \leq \text{value}$  in node;
        result ← deleteValue (childID, value);
        if result = true then // nothing to do
            | return true
        else // the child has become under-full
            child ← get node with ID childID;
            if child is the only child of node then // node must be the root
                | delete the old root node;
                | set child as the new root node;
                | return true
            else // child has a neighbor below node
                neighbor ← get a neighbor of child;
                if neighbor has more than  $d$  values then
                    | re-fill child by borrowing from neighbor;
                    | return true
                else // neighbor is minimally full
                    | merge child with neighbor;
                    | adjust the pointers and values of node;
                    | delete child;
                    | if node still has enough values stored then
                        | | return true
                    | else // node is under-full
                        | | return false

```

**Algorithm 3:** Deleting a value from a B+ Tree.

This page intentionally left blank.



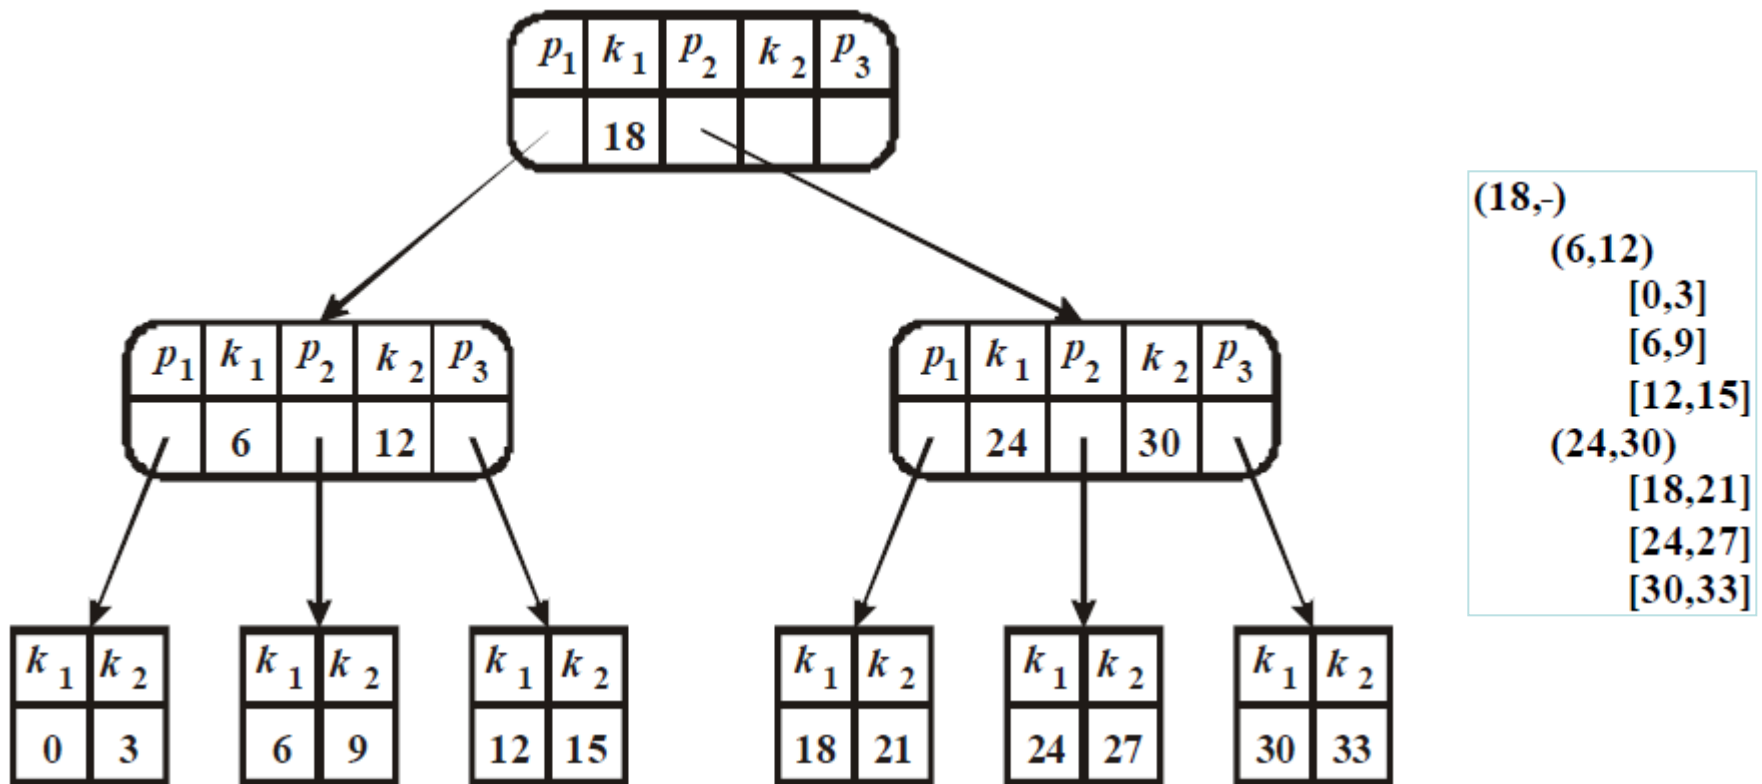


# HW #5 B+ Tree

- You will implement a simpler version (**insert** only) of B+ Tree using **C** / **C++** / **Java** / **Python**, or **other** programming language if you prefer.
- Your tree will be created and maintained in **memory** rather than on **disk** as is commonly the case.
- The task is to write a program that reads integers, store them into a B+ Tree, prints the tree, and prints a sorted list of identifiers.

# HW #5 (2)

- A B+ Tree can be represented by listing the nodes using a **preorder** traversal and indenting them according to their depth.



# HW #5 (3)

- The program (BTree) is invoked with command line arguments of **variable** numbers of program parameters. The first program parameter is the **key capacity**; then followed by a list of positive integers.
- Sample command lines together with a representative list of integers are shown below:

```
$ BTree 3 22 16 41 11 18 28 58 1 8 12 17 19 23 31 52 59 61
```

# HW #5 (4)

- Program **BTree** reads an integer key, inserts it into a B+ Tree, and prints a representation of the tree after each integer is inserted. Note that the integer about to be inserted is printed just before the B+ Tree is printed. B+ Trees are printed using a preorder traversal.
- Keys in interior nodes of the B+ Tree are enclosed in parentheses and separated by colons (:). Empty positions for keys in a node that is not completely filled are represented by **underscores** (or **dash** if you prefer).
- Nodes in the B+ Tree are indented according to the **depth** of the node. The root is not indented. Each level is indented **five** spaces more than the previous level.
- Following shows the output produced by the program given in the sample list of integers.

```

Insert key 22:
[22,_,_]

Insert key 16:
[16,22,_]

Insert key 41:
[16,22,41]

Insert key 11:
(22:_)
  [11,16,_]
  [22,41,_]

Insert key 18:
(22:_)
  [11,16,18]
  [22,41,_]

Insert key 28:
(22:_)
  [11,16,18]
  [22,28,41]

Insert key 58:
(22:41)
  [11,16,18]
  [22,28,_]
  [41,58,_]

Insert key 1:
(22:_)
  (16:_)
    [1,11,_]
    [16,18,_]
  (41:_)
    [22,28,_]
    [41,58,_]

```

```

Insert key 8:
(22:_)
  (16:_)
    [1,8,11]
    [16,18,_]
  (41:_)
    [22,28,_]
    [41,58,_]

Insert key 12:
(22:_)
  (11:16)
    [1,8,_]
    [11,12,_]
    [16,18,_]
  (41:_)
    [22,28,_]
    [41,58,_]

Insert key 17:
(22:_)
  (11:16)
    [1,8,_]
    [11,12,_]
    [16,17,18]
  (41:_)
    [22,28,_]
    [41,58,_]

```

```

Insert key 19:
(16:22)
  (11:_)
    [1,8,_]
    [11,12,_]
  (18:_)
    [16,17,_]
    [18,19,_]
  (41:_)
    [22,28,_]
    [41,58,_]

Insert key 23:
(16:22)
  (11:_)
    [1,8,_]
    [11,12,_]
  (18:_)
    [16,17,_]
    [18,19,_]
  (41:_)
    [22,23,28]
    [41,58,_]

```

```

Insert key 23:
(16:22)
  (11:_)
    [1,8,_]
    [11,12,_]
  (18:_)
    [16,17,_]
    [18,19,_]
  (41:_)
    [22,23,28]
    [41,58,_]

```

```

Insert key 31:
(16:22)
  (11:_)
    [1,8,_]
    [11,12,_]
  (18:_)
    [16,17,_]
    [18,19,_]
  (28:41)
    [22,23,_]
    [28,31,_]
    [41,58,_]

```

```

Insert key 52:
(16:22)
  (11:_)
    [1,8,_]
    [11,12,_]
  (18:_)
    [16,17,_]
    [18,19,_]
  (28:41)
    [22,23,_]
    [28,31,_]
    [41,52,58]

```

```

Insert key 59:
(22:_)
  (16:_)
    (11:_)
      [1,8,_]
      [11,12,_]
    (18:_)
      [16,17,_]
      [18,19,_]
  (41:_)
    (28:_)
      [22,23,_]
      [28,31,_]
  (58:_)
    [41,52,_]
    [58,59,_]

```

```

Insert key 61:
(22:_)
  (16:_)
    (11:_)
      [1,8,_]
      [11,12,_]
    (18:_)
      [16,17,_]
      [18,19,_]
  (41:_)
    (28:_)
      [22,23,_]
      [28,31,_]
  (58:_)
    [41,52,_]
    [58,59,61]

```

# HW #5 (5)

## Second test case:

- Following shows with another list of integers with corresponding actions (and outputs):

\$ BTree 2 51 29 73 105 15 31

1. Start with an empty B-tree of order 1. Insert key **51**. The resulting tree contains a single leaf.

**[51,-]**

2. Insert key **29**. Key 51 is shifted one position to the right preserving the property that keys are always in ascending order.

**[29,51]**

3. Insert key **73**. Key 73 is appended to the list of keys in the node. The leaf is now *overflow*.

**[29,51,73]**

3.1. Split the leaf.

**[29,-]**

**[51,73]**

3.2. Create an empty interior node that will be the parent of the two leaves.

**(:-)**

**[29,-]**

**[51,73]**

3.3. *Copy* the smallest key in the node having the larger keys into the newly created interior node.

**(51:-)**

**[29,-]**

**[51,73]**

3.4. Assign a pointer to the first leaf **[29,-]** to  $p_1$  and a pointer to the second leaf, **[51,73]** to  $p_2$ .



4. Insert key 105.

(51:-)

[29,-]

[51,73,105]

4.1. The leaf, [51,73,105], is *overfull*.

4.2. Split the leaf

(51:-)

[29,-]

[51, -]

[73,105]

4.3. Create an empty interior node that will be the parent of the two leaves.

(51:-)

[29,-]

(-, -)

[51, -]

[73,105]

4.4. *Copy* the smallest key in the node having the larger keys into the newly created interior node.

**(51:-)**

**[29,-]**

**(73,-)**

**[51, -]**

**[73,105]**

4.5. Merge the newly created interior node with the existing interior node.

**(51:73)**

**[29,-]**

**[51, -]**

**[73,105]**

5. Insert key **15**.

**(51:73)**

**[15,29]**

**[51, -]**

**[73,105]**

6. Insert key 31.

(51:73)

[15,29,31]

[51, -]

[73,105]

6.1. The leaf [15,29,31], is *overflow*. Split the leaf into [15,-] and [29,31]. Hoist the middle key into a new interior node (29:-) having pointers to the two leaves [15,-] and [29,31].

(51:73)

(29:-)

[15,-]

[29,31]

[51, -]

[73,105]

6.2. The new interior node must be merged with the existing interior node.

(51:73) ? (29:-)

[15,-]

[29,31]

[51, -]

[73,105]

(29:51:73)

[15,-]

[29,31]

[51, -]

[73,105]

6.3. The interior node, (29:51:73) is *overflow* and must be split.

(29:51:73)

[15,-]

[29,31]

[51, -]

[73,105]

6.4. Splitting an interior node is different than splitting a leaf. The middle key is *removed* rather than copied. Pointers to the subordinate interior nodes are assigned to pointers,  $p_1$  and  $p_2$  on either side of key **51** in the new root (**51:-**).

(**51:-**)

(**29:-**)

(**73:-**)

6.5. Subordinate leaves remain attached to their respective interior nodes.

(**51:-**)

(**29:-**)

[15,-]

[29,31]

(**73:-**)

[51, -]

[73,105]