

作業十：追蹤 execve 系統呼叫

學習目標：

1. 了解作業系統如何解析執行檔案的格式
2. 了解作業系統在處理 execve 時，是否立即載入執行檔案？或者只是修改 task_struct 中的 mm_struct？

題目：

1. 撰寫程式碼稱之為 myls，在程式碼中使用 execve 系列的任何 libc 函數，載入新的執行檔案 (ls)。注意：不需要使用 fork。
2. 請問作業系統如何載入執行檔案？附上程式碼的截圖，並約略說明。
3. 請問作業系統是否立即載入執行檔案到記憶體中？附上程式碼的截圖並約略說明

● 作業繳交：

- 上述 1~3 都是繳交 pdf。其中(1)必須附上你的程式碼的截圖。
- Pdf 中的學號、姓名（請隱藏個人資訊，例如：學號 687410007，姓名：羅 X 五）

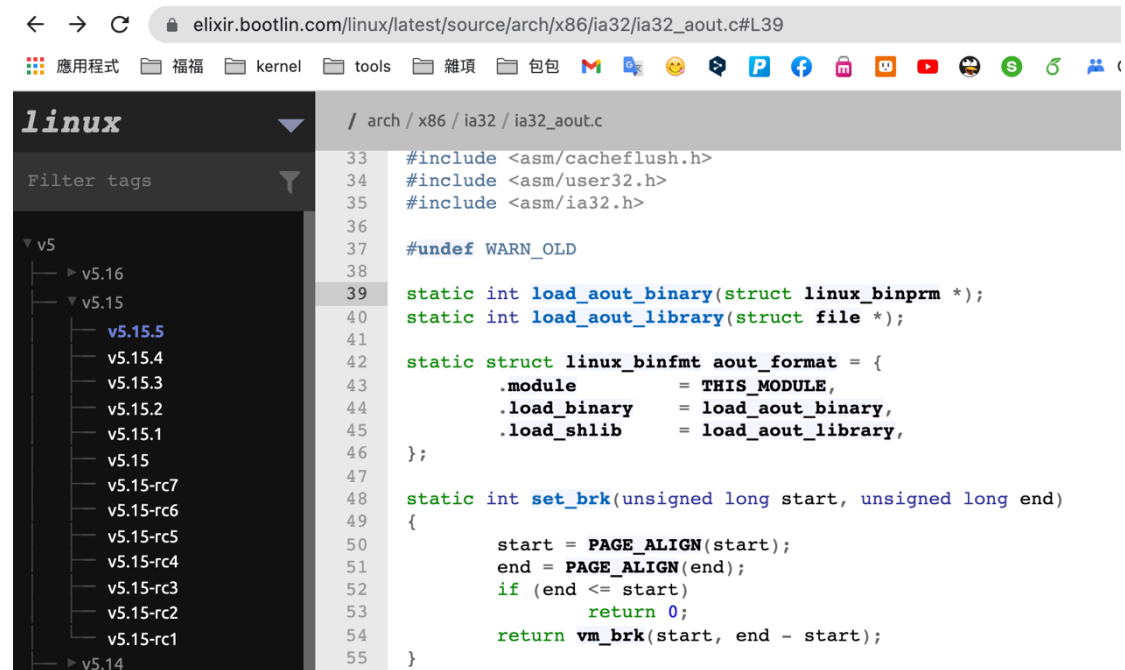
繳交：

- 繳交期限：請參考網頁
- **這次沒有 youtube 的說明。**如果真的不會寫，記得去請教朋友。在你的報告上寫你請教了誰即可。

提示:

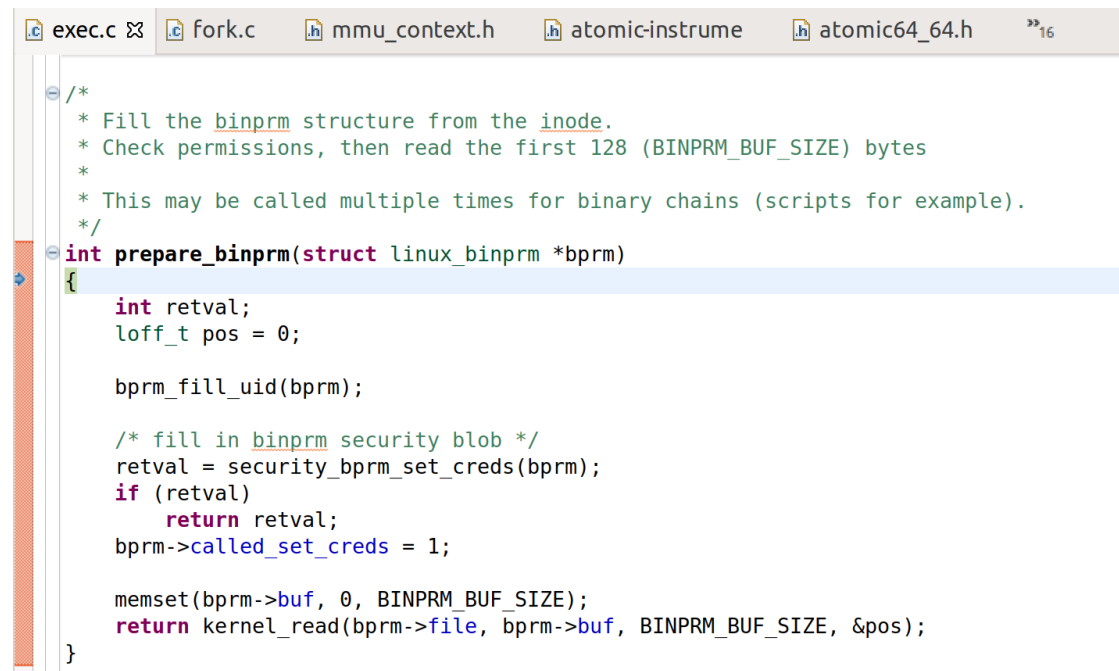
1. 可以將中斷點設定在 b do_execve

2. 圖如下



```
← → ↻ elixir.bootlin.com/linux/latest/source/arch/x86/ia32/ia32_aout.c#L39
應用程式 福福 kernel tools 雜項 包包
linux / arch / x86 / ia32 / ia32_aout.c
Filter tags
v5
  v5.16
  v5.15
    v5.15.5
    v5.15.4
    v5.15.3
    v5.15.2
    v5.15.1
    v5.15
    v5.15-rc7
    v5.15-rc6
    v5.15-rc5
    v5.15-rc4
    v5.15-rc3
    v5.15-rc2
    v5.15-rc1
  v5.14
33 #include <asm/cacheflush.h>
34 #include <asm/user32.h>
35 #include <asm/ia32.h>
36
37 #undef WARN_OLD
38
39 static int load_aout_binary(struct linux_binprm *);
40 static int load_aout_library(struct file *);
41
42 static struct linux_binfmt aout_format = {
43     .module      = THIS_MODULE,
44     .load_binary  = load_aout_binary,
45     .load_shlib   = load_aout_library,
46 };
47
48 static int set_brk(unsigned long start, unsigned long end)
49 {
50     start = PAGE_ALIGN(start);
51     end = PAGE_ALIGN(end);
52     if (end <= start)
53         return 0;
54     return vm_brk(start, end - start);
55 }
```

3. 圖如下



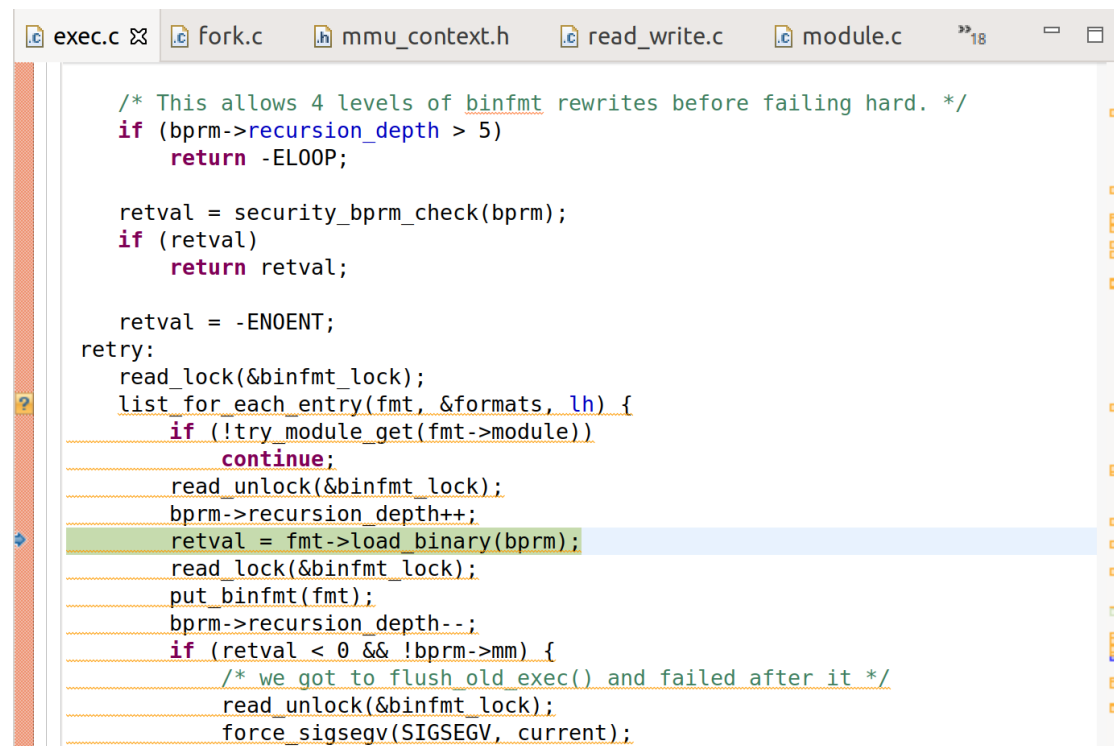
```
exec.c  fork.c  mmu_context.h  atomic-instrume  atomic64_64.h  »16
/*
 * Fill the binprm structure from the inode.
 * Check permissions, then read the first 128 (BINPRM_BUF_SIZE) bytes
 *
 * This may be called multiple times for binary chains (scripts for example).
 */
int prepare_binprm(struct linux_binprm *bprm)
{
    int retval;
    loff_t pos = 0;

    bprm_fill_uid(bprm);

    /* fill in binprm security blob */
    retval = security_bprm_set_creds(bprm);
    if (retval)
        return retval;
    bprm->called_set_creds = 1;

    memset(bprm->buf, 0, BINPRM_BUF_SIZE);
    return kernel_read(bprm->file, bprm->buf, BINPRM_BUF_SIZE, &pos);
}
```

4. 圖如下



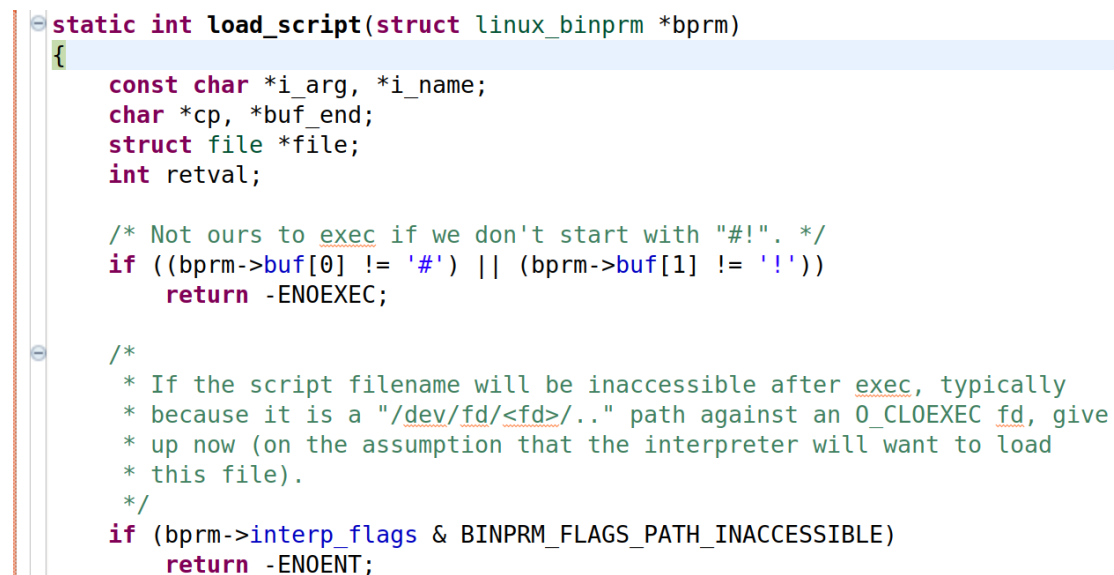
```
exec.c  fork.c  mmu_context.h  read_write.c  module.c  18

/* This allows 4 levels of binfmt rewrites before failing hard. */
if (bprm->recursion_depth > 5)
    return -ELOOP;

retval = security_bprm_check(bprm);
if (retval)
    return retval;

retval = -ENOENT;
retry:
    read_lock(&binfmt_lock);
    list_for_each_entry(fmt, &formats, lh) {
        if (!try_module_get(fmt->module))
            continue;
        read_unlock(&binfmt_lock);
        bprm->recursion_depth++;
        retval = fmt->load_binary(bprm);
        read_lock(&binfmt_lock);
        put_binfmt(fmt);
        bprm->recursion_depth--;
        if (retval < 0 && !bprm->mm) {
            /* we got to flush_old_exec() and failed after it */
            read_unlock(&binfmt_lock);
            force_sigsegv(SIGSEGV, current);
        }
    }
}
```

5. 圖如下



```
static int load_script(struct linux_binprm *bprm)
{
    const char *i_arg, *i_name;
    char *cp, *buf_end;
    struct file *file;
    int retval;

    /* Not ours to exec if we don't start with "#!". */
    if ((bprm->buf[0] != '#') || (bprm->buf[1] != '!'))
        return -ENOEXEC;

    /*
     * If the script filename will be inaccessible after exec, typically
     * because it is a "/dev/fd/<fd>/" path against an O_CLOEXEC fd, give
     * up now (on the assumption that the interpreter will want to load
     * this file).
     */
    if (bprm->interp_flags & BINPRM_FLAGS_PATH_INACCESSIBLE)
        return -ENOENT;
}
```

6. 圖如下

```
static int load_elf_binary(struct linux_binprm *bprm)
{
    struct file *interpreter = NULL; /* to shut gcc up */
    unsigned long load_addr = 0, load_bias = 0;
    int load_addr_set = 0;
    char * elf_interpreter = NULL;
    unsigned long error;
    struct elf_phdr *elf_ppnt, *elf_phdata, *interp_elf_phdata = NULL;
    unsigned long elf_bss, elf_brk;
    int bss_prot = 0;
    int retval, i;
    unsigned long elf_entry;
    unsigned long interp_load_addr = 0;
    unsigned long start_code, end_code, start_data, end_data;
    unsigned long reloc_func_desc __maybe_unused = 0;
    int executable_stack = EXSTACK_DEFAULT;
    struct pt_regs *regs = current_pt_regs();
    struct {
        struct elfhdr elf_ex;
```

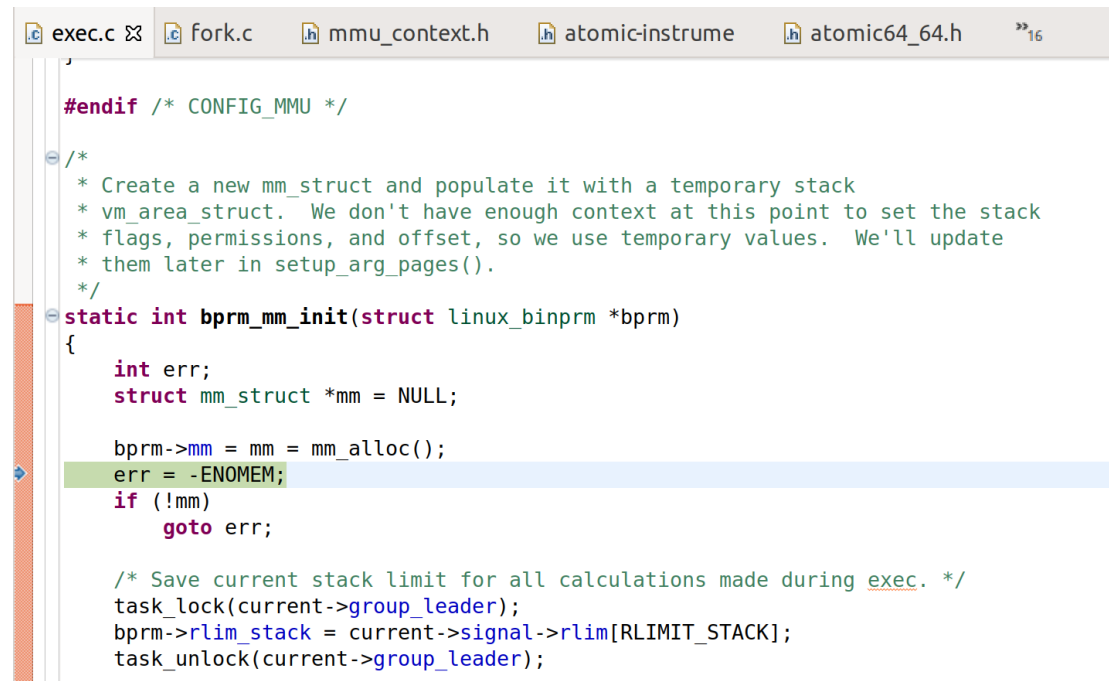
7. 圖如下

The screenshot shows a GDB (9.2) interface with the 'Project Explorer' and 'Debug' tabs. The 'Debug' tab is active, showing a running process 'RunningLinuxKernel 5.0 [C/C++ Remote Application]'. The call stack is expanded, showing the current frame 'load_elf_binary() at binfmt_elf.c:691 0xffffffff81538ea4'. The stack trace includes the following frames (from top to bottom):

- load_elf_binary() at binfmt_elf.c:691 0xffffffff81538ea4
- search_binary_handler() at exec.c:1,656 0xffffffff81484b27
- exec_binprm() at exec.c:1,698 0xffffffff81484d9d
- __do_execve_file() at exec.c:1,818 0xffffffff814854b1
- do_execveat_common() at exec.c:1,865 0xffffffff81485737
- do_execve() at exec.c:1,882 0xffffffff81485838
- __do_sys_execve() at exec.c:1,963 0xffffffff81485c94
- __se_sys_execve() at exec.c:1,958 0xffffffff81485c4f
- __x64_sys_execve() at exec.c:1,958 0xffffffff81485be9
- do_syscall_64() at common.c:290 0xffffffff810075e6
- <...more frames...>

The GDB version is 9.2.

8. 圖如下



```
exec.c  fork.c  mmu_context.h  atomic-instrume  atomic64_64.h  16

#endif /* CONFIG_MMU */

/*
 * Create a new mm_struct and populate it with a temporary stack
 * vm_area_struct. We don't have enough context at this point to set the stack
 * flags, permissions, and offset, so we use temporary values. We'll update
 * them later in setup_arg_pages().
 */
static int bprm_mm_init(struct linux_binprm *bprm)
{
    int err;
    struct mm_struct *mm = NULL;

    bprm->mm = mm = mm_alloc();
    err = -ENOMEM;
    if (!mm)
        goto err;

    /* Save current stack limit for all calculations made during exec. */
    task_lock(current->group_leader);
    bprm->rlim_stack = current->signal->rlim[RLIMIT_STACK];
    task_unlock(current->group_leader);
}
```